



Fall 2019

IoT Data Analytics



Vassilis Christophides

christop@csd.uoc.gr

<http://www.csd.uoc.gr/~hy562>

University of Crete, Fall 2019

Unia



Fall 2019

The Internet of Things (IoT)

- Networks of **physical objects** (aka Things) that are **uniquely identifiable** in the Internet with embedded **sensing** and **actuating** along with **programmability** capabilities
- Information **about** Things can be collected and the **state** of Things can be changed from **anywhere**, **anytime**, by **anything** !



Unia

<http://www.internet-of-things-research.eu>

1



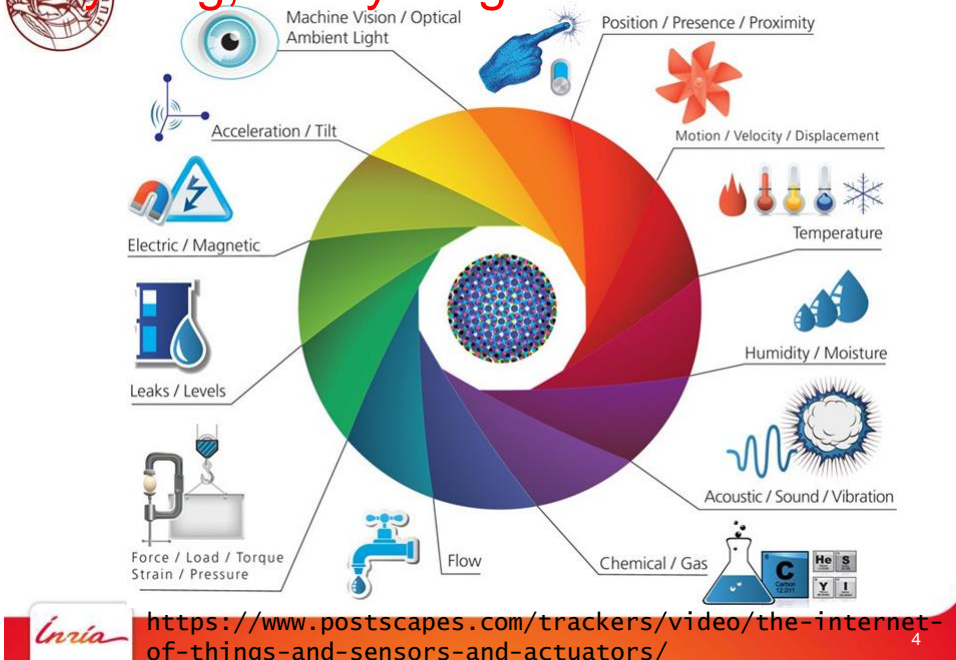
Fall 2019

A World of Things



Fall 2019

Anything, Everything can be Measured!



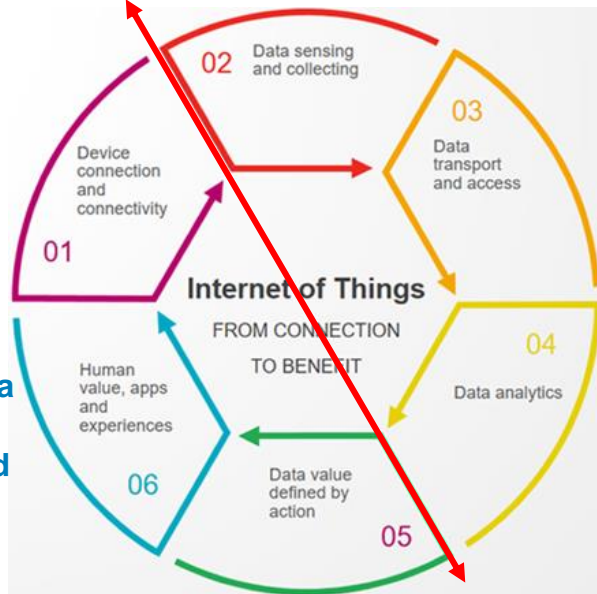
<https://www.postscapes.com/trackers/video/the-internet-of-things-and-sensors-and-actuators/>



Fall 2019

The IoT Value Chain

- IoT has the potential to transform *how* and *when* decisions are made throughout business and our daily lives iff *high quality data* can be *processed efficiently* and *analyzed effectively*!



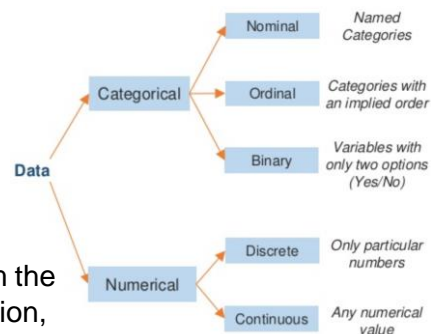
<https://www.besanttechnologies.com/iot-training-in-bangalore5>



Fall 2019

Types of IoT Data

- Addresses/Unique Identifiers** (nominal)
 - IP (IPv4, IPv6), Bluetooth, Zigbee, LoRa, RFID
- Positional Data** (continuous)
 - GPS (Longitude, Latitude, Altitude), WiFi (Longitude, Latitude)
- Temporal Data** (continuous)
 - Time and Date
- Sensor Data** (numerical)
 - Output of a device that detects and responds to some type of input from the physical environment (motion, position, environment, mass, biomarker)
- Descriptive Data** (categorical)
 - Objects, Processes, and Systems



7



IoT Data: Sensing the Physical World!

- IoT data is becoming soon “mega” **big** (i.e. billions of connected Things)
 - 44EB / months corresponds to ~ 100M hard disks
 - Processing 44 EB with 100M servers would still take > one hour
- IoT “**data in motion**” as opposed to traditional “data at rest”
 - **Data speed** is bound by the sensing frequency and connectivity
 - **Data throughput** (vs processing) **has become the limiting factor!**
- **High variety** data e.g., from *heterogeneous Things, embedded in the environment or wearable by persons*
- **High veracity** data (dropped or unlivable) as IoT data quality depends on *how Things are used and connected in the wild*
- **High variability** data as Things *frequently change behavior dynamically* and in ways that are not fully known in advance
- Data **security**, **privacy**, etc. even more important!



IoT Data Analytics

- **Use Cases**
 - Real-time monitoring and control
 - Failure detection and predictive maintenance **Anomaly Detection**
 - Malicious cyber activity detection
 - Product life-cycle management
 - Operational efficiency, optimization, self-adaptation
- **Batch and Online Techniques**
 - Descriptive Analytics: aggregation and statistics
 - Predictive Analytics: Supervised, **Semi-Supervised, Unsupervised**
 - Data series analytics: Spatiotemporal
 - Domain and data type specific: Signal processing



Real Time Anomaly Detection

Fall 2019



Monitor Traffic Conditions



Network and Cyber Security
Detect Intrusion



Monitor Biomarkers



Predictive maintenance,
Production safety
Monitoring



Agriculture
Pest, Water
Control



Monitor Energy
Consumption

Define **alerts**, predict **faults**, detect **intrusions** & **threats**

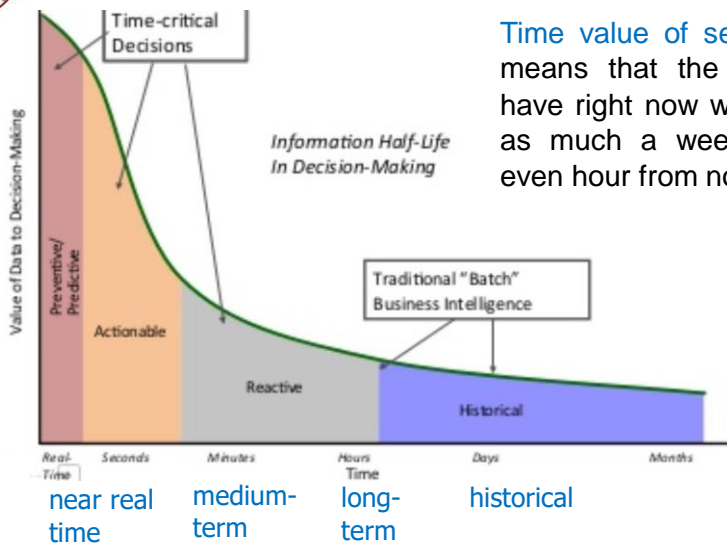
Inria

11



What is Real Time Analytics?

Fall 2019



Time value of sensor data means that the data you have right now won't mean as much a week, day or even hour from now !

Inria

<https://www.slideshare.net/AmazonWebServices/introduction-to-realtime-streaming-data-and-amazon-kinesis-streaming-data-ingestion-with-firehouse>

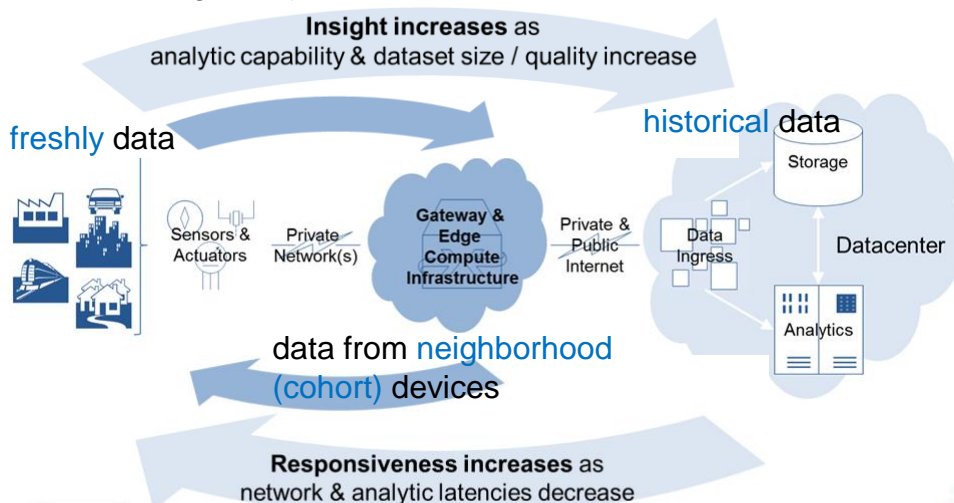
13



Fall 2019

A 3-tier Architecture for IoT Analytics

- To support analytics with different time horizons (near-real time, medium, long terms), we need to a **3-tier architecture** for IoT



<http://pjtec.info/how-the-internet-of-things-will-shape-the-datacenter-of-the-future-2/iot-20150804/>

14



Fall 2019

Edge Processing/Computing

- Objective:** push data processing away from the core and towards the edge of the network
 - help ensuring that the **right data processing task** takes place at the **right time** and **place**
- Motivating factors:**
 - Reduce latency:** run data computations directly on IoT devices or gateways, and only interact with the Cloud off the critical path (e.g. to continuously train ML models with freshly data)
 - Be robust to connectivity issues:** applications are not disrupted in case of limited or intermittent network connectivity
 - Preserve privacy:** ensures that sensitive data is pre-processed **on-site**, and **only data that is privacy compliant** is sent to the Cloud for further analysis, after having passed through a first layer of anonymizing aggregation
- Wide range of technologies:** Local Cloud/Fog computing, Grid/Mesh Computing, mobile edge computing, cloudlets, etc.

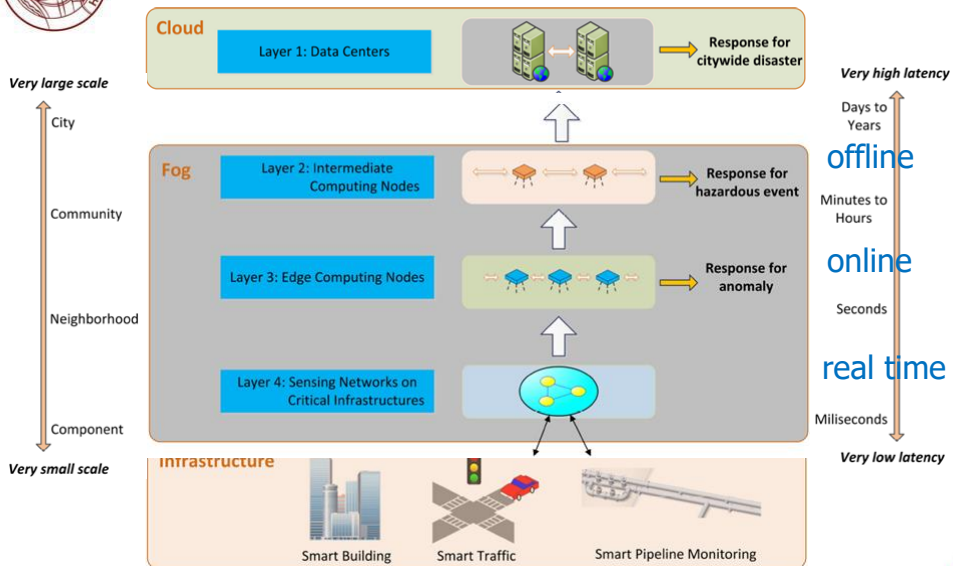


15



Fall 2019

A Fog/Edge Computing Architecture



B. Tang, Z. Chen, G. Heffernan, T. Wei, H. He, Q. Yang. A Hierarchical Distributed Fog Computing Architecture for Big Data Analysis in Smart Cities. ASE BD&SI 2015

16



Fall 2019

So to Recap


- The Internet of Things Is more About Data, than Things!
 - data in motion usually with spatiotemporal semantics
 - produced by several, distributed, heterogeneous devices
 - in the wild, hence data prone to errors and noise
- Low-latency, privacy-aware IoT data processing in post-cloud architectures
 - Edge/fog computing for effectively and efficiently ingesting and analyzing data when and where it is needed
- Real time IoT Data Analytics requires data streams support
 - Limited ground truth calls for unsupervised techniques



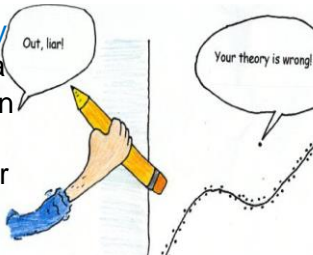
17




Data Anomaly Detection



Data Anomalies

- What are anomalies?
 - “An outlier is an observation in a dataset that appears to be inconsistent with the remainder of that dataset” [Johnson 1992]
 - “An outlier is an observation that deviates so much from other observations as to arouse suspicion that it was generated by a different mechanism” [Hawkins 1980]
 - Anomaly, or outlier, or deviation, or novelty detection, aims to measure whether a data point (or a subset) considerably differs than the remainder of the data points
 - On a scatter plot of the data, they lie far away from other data
 - Anomalies = Outliers + Novelties
- 
- A cartoon illustration of a hand holding a yellow pencil, pointing at a scatter plot. The hand is wearing a blue sleeve. Two speech bubbles are present: one on the left saying "Out, liar!" and one on the right saying "Your theory is wrong!". The scatter plot shows a cluster of points with a few points far away from the cluster.






Cause of Data Anomalies

Fall 2019

- Errors in the data collection or measurement process
 - Because of human **error** (intentional or not), a **problem** with a measuring device or the presence of noise
 - Insights extracted from “dirty data” are probably erroneous and thus the decisions to be made are likely **unsound** (e.g., incur a high rate of *false positives* and *false negatives*)
- Changes in the “data generation” process, e.g. some given statistical process (not an error, **novelties** in data)
 - Abnormal data deviate from this generating mechanism because it is of a **different type**
 - Novel class of observations provide valuable insights**



inria


20



Point Anomaly Detection


Types of Anomalies

Fall 2019




Data point anomalous w.r.t. to the rest of the data

Contextual Anomaly Detection



Data point anomalous in a specific **context**; also referred to as **conditional anomalies**

Group Anomaly Detection



A **group** of data points is anomalous; the individual points within a group are not anomalous by themselves

inria

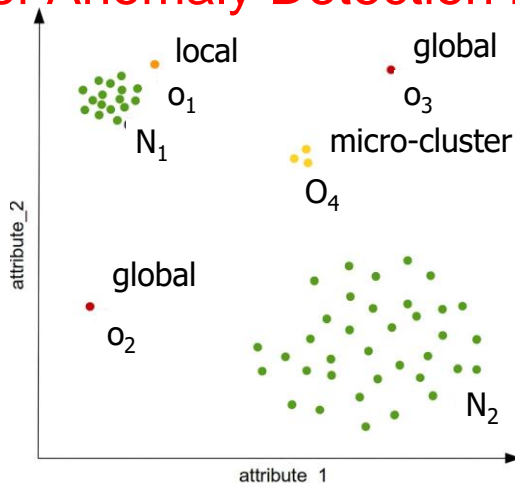
Chandola, V., Banerjee, A., & Kumar, V. (2009). Anomaly detection: A survey. ACM Computing Surveys. 41(3). 1–58.

21



Fall 2019

Scope of Anomaly Detection Methods



- **Swamping**: wrongly identifying normal instances as anomalies when normal instances are too close to anomalies
- **Masking**: an anomaly is close to another point, anomalous or not



M. Goldstein, S. Uchida A Comparative Evaluation of Unsupervised Anomaly Detection Algorithms for Multivariate Data. PLoS One. 2016:11(4)

22



Fall 2019

Example: Most Unusual Objects?




<https://www.slideshare.net/mlv1c/114-anomaly-detection>

23



Example: How to Distinguish Objects?




High-Dimensional Data: Example

		Object	Length/Width	Num Surfaces	Smooth	round
		penny	1	3	true	
		dime	1	3	true	
		knob	1	4	true	
corners		box	1	6	true	skinny
		eraser	2.75	6	true	
		block	1.6	6	true	
		screw	8	3	true	
		battery	5	3	true	
		key	4.25	3	false	
		bead	1	2	true	



Fall 2019

Multi-Dimensional Anomaly Detection



- Anomaly detection in **multi-dimensional** datasets reveals more accurate behavior of the data but at the same time poses various challenges




<http://wikistat.fr/pdf/st-m-app-anomalies.pdf>

27



Fall 2019

Multi-Dimensional Anomaly Detection



- Space concentration of distances**: feature-wise distances of i.i.d. data samples approximately converge to a normal distribution
- The number of feature subspaces grows exponentially** with the increasing dimensionality of data: **data density functions** not easily computed
- Data-snooping bias**: given enough features, at least one feature subspace can be found for each data point such that it appears as an anomaly





D. L. Donoho, "High-dimensional data analysis: The curses and blessings of dimensionality," AMS Math Challenges Lecture, pp. 1–32, 2000.

28



Multi-Dimensional Anomaly Detection



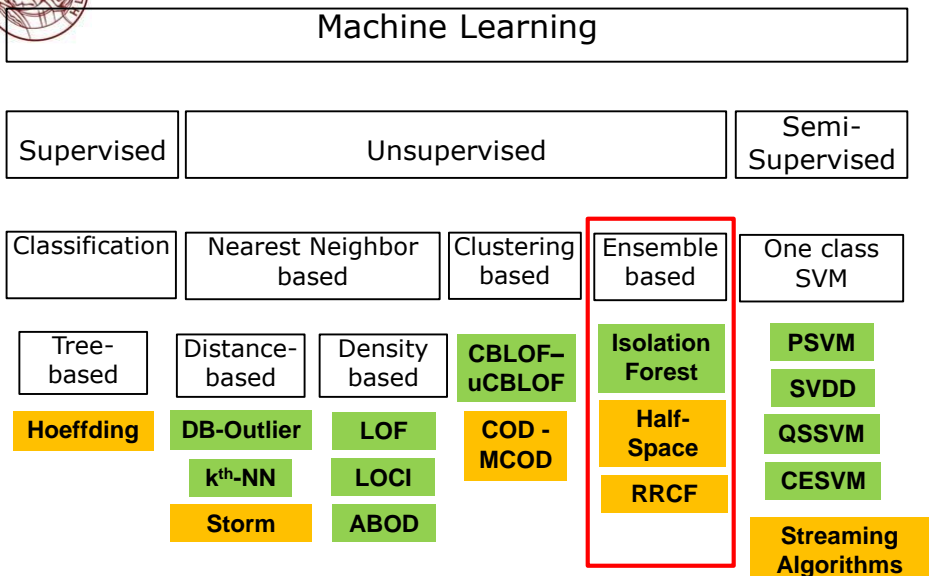
- **Challenge:** select a subspace that highlights *relevant features*, i.e. in which anomalies exhibit significantly different values from normal data
 - A *robust* anomaly detector should be able to detect anomalies with a *high proportion of irrelevant features creating noise in the input data, which masks the true anomalies*



L. Wilkinson Visualizing Big Data Outliers through Distributed Aggregation
IEEE Trans. Vis. Comput. Graph. 24(1): Jan. 2018



Anomaly Detection Methods




Goldstein M, Uchida S (2016) A Comparative Evaluation of Unsupervised Anomaly Detection Algorithms for Multivariate Data. PLOS ONE 11(4)



Random Trees

Fall 2019

Build a decision tree from a random data subsample



- Partition feature space using axis-parallel subdivisions
 - Select the split feature A , randomly and uniformly
 - Select the split value V_A , uniformly as the $\min(A) + (\max(A) - \min(A)) * \text{rand}(1)$
- Grow a random tree until each data point is in its own leaf or the tree reaches a maximum height

Unia


<https://www.slideshare.net/mlv1c/114-anomaly-detection>



Isolation Forest [Liu et al. 2008]

Fall 2019

- To score a data point, find the height of the leaf node
 - The smaller the height the more anomalous is the data



- Build an ensemble of decision trees from randomly selected subsamples of size n


Unia

<https://www.slideshare.net/mlv1c/114-anomaly-detection>



Isolation Forest Scores

- Use **average height** to compute the **anomaly score**:
 - 0 (normality) 1 (abnormality)



Score

- Ensemble average path length to a data point
- Normalized by the **expected path length** of balanced binary search tree

$$s(x, n) = 2^{-\frac{E(h(x))}{c(n)}}$$




<https://towardsdatascience.com/a-brief-overview-of-outlier-detection-techniques-1e0b2c19e561>

35



Isolation Forest Scores



(a) Isolating x_i

(b) Isolating x_o

12 partitions (**not an anomaly**)

4 partitions (**anomaly**)



<https://www.depends-on-the-definition.com/detecting-network-attacks-with-isolation-forests/>

36



iForest: Pros and Cons

- Pros
 - Very **easy to construct** (no distance/density function needed) avoiding hard decisions whether a data point is an anomaly or not
 - assigns an anomalous score to each data point
 - Achieve a **sublinear time-complexity** and a **small memory-footprint**
 - By exploiting **subsampling**
 - By eliminating major computational cost of distance calculation in all the distance-based and density-based AD methods
 - Can provide **anomaly explanations** [Siddiqui et al. 2015]
- Cons
 - Hyper-parameter tuning (e.g. number/height of trees, sample size)
 - Large datasets will need more isolation trees (how many?)
 - Requires a **high percentage of relevant features to identify anomalies** [Bandaragoda et al. 2018]
 - In presence of features that do not provide information over the anomaly, iForest increases height randomly by ignoring this fact



Performance Measures of AD Algorithms


Confusion Matrix	Actual Normal Data (n_n)	Actual Anomalous Data (n_a)
Predicted Non-anomalies	TN	FN
Predicted Anomalies	FP	TP

- **Accuracy Rate (ACC)** = $(TP + TN) / (TN + FP + FN + TP)$
- **False Alarm –Positive– Rate (FAR)** = $FP / (FP + TN)$
- **True Positive –Detection– Rate (TPR)** = $TP / (TP + FN)$
- **Receiver Operating Characteristic (ROC)**= tradeoff between **TPR&FAR**
- **Area Under the ROC Curve (AUC)**: can be computed by a slight modification of the algorithm for constructing **ROC** curves



ROC Spaces and AUC

- ROC curves are two dimensional plots in which the **true positive (TP) rate** is plotted on the Y axis and the **false positive (FP) rate** on the X axis
- Area Under the ROC Curve (AUC) has an important statistical property:
 - The AUC of a classifier is equivalent with the **probability that the classifier will rank a randomly chosen positive instance higher than a randomly chosen negative instance**
 - When comparing classifiers, **the bigger AUC the better!**




Summary

- AD algorithms **usually report different outliers per dataset**
 - they may miss obvious outliers especially in heterogeneous datasets that contain a mixture of numerical and categorical attributes
- Which anomaly detection method to use depends on
 - Data characteristics:** *dimensionality* (Univariate, Multivariate) and *type* (categorical numerical)
 - Anomaly characteristics:** *type* (Point, Collective), *semantics* (Set-based, Sequence-based), *kind* (Binary, Score)
 - Availability of labels:** *supervised*, *unsupervised*, *semi-supervised*
 - Algorithmic properties:** *computational cost* (exponential vs linear), *prior knowledge* (parametric, non-parametric), *distributional/incremental computation potential*, ...
- How we can **automate the selection of the most suitable AD algorithm for a particular dataset** with minimal human involvement and within limited computational budgets?




Anomaly Detection in Data Streams



Offline vs Online Anomaly Detection

- A **data stream** is a possible infinite series of data points $\dots, o_{n-2}, o_{n-1}, o_n, \dots$, where data point o_n is received at time $o_n.t$.




- **Offline learning** (for finite streams):
 - All data is **stored** and can be accessed in **multiple passes**
 - Learned **model is static**
 - Potentially **high memory and processing overhead**
- **Online learning** (for infinite streams):
 - **Single pass** over the data which are discarded after being processed
 - Learned **model is updated** : one point at a time or in mini-batches
 - **Low memory and processing overhead**




Online Anomaly Detection: Challenges

- Limited computational resources for unbounded data
 - window-based algorithms
- High speed data streams:
 - the rate of updating a detection model should be higher than the arrival rate of data
- Both normal and anomalous data characteristics may evolve in time and hence detector models need to quickly adjust to concept drifts/shifts:
 - adversarial situations (fraud, insider threats)
 - diverse set of potential causes (novel device failure modes)
 - user's notion of “anomaly” may change with time



Robust Random Cut Forest (RRCF) [Guha et al, ICML 2016]


- Construct random trees as iForest but incrementally update their binary tree structure using sliding windows
 - Traverse a tree to Insert a new point to a leaf node, use FIFO to remove the oldest point if this procedure exceeds the max number of leaf nodes
 - Calculate the Displacement of the new point, as the number of leaf nodes that updated (moved/deleted) in a tree after the insertion
- Principled definitions of anomalies
 - Calculate the Collusive Displacement (CD) of a new point in a tree, as the maximal displacement across all nodes traversed in a path from the root
 - Outliers correspond to points that their average CD in the forest is large, instead of small average tree height
- Address iForest limitation in case of multiple irrelevant features
 - Does not uniformly selects the split feature





RRCF : Example

- Given the following input stream of points S



- Start by collecting the **first** $w = 7$ points of S , as a **Initial Training Sample** $x = \{x_1, x_2, \dots, x_j\}$ where $x_i \in R^d$, in order to build the initial Tree
- Assumption :
 - all objects within x are considered as **inliers**



RRCF Tree Construction: Example

- Consider the **value range** per feature

$$l_i = \max_x x_i - \min_x x_i$$

- Pick a **splitting feature** w.r.t. its normalized value range $\frac{l_i}{\sum_j l_j}$

- Choose uniformly a **splitting value**
 $X_i \sim \text{Uniform}[\min_x x_i, \max_x x_i]$


Object	L/W	NS	S
Screw	8	3	0
Dime	1	3	1
Knob	1	4	1
Battery	5	3	1
Box	1	6	1
Block	1.6	6	1
Penny	1	2	1

	L/W	NS	S
Minimum value	1	2	0
Maximum Value	8	6	1
Range	7	4	1



Fall 2019

RRCF Tree Construction: Example



Object	L/W	NS	S
Dime	1	3	1
Knob	1	4	1
Penny	1	2	1
Range	0	2	0
Penny	1	2	1
Range	0.6	4	0

- Note that: Each (internal) child node has a sub-space (bound) feature value range of its parent node



Length/Width (L/W), Num Surface (NS), Smooth (S)


56



Fall 2019

RRCF Tree Update: Example

- Consider the new object *Eraser*, by sliding the window by one point



Object	L/W	NS	S
Eraser	2.75	6	1


	L/W	NS	S
Minimum value	1	6	1
Maximum Value	1.6	6	1
Value			

- A new tree instance has been computed
 - In theory update tree to the new instance only if low $\text{CoDisp}(\text{Eraser})$





RRCF Tree Update: Example



$$\text{CoDisp}(\text{Eraser}_4) = \frac{\text{Data}(\text{Displacement})}{\text{leaves of sibling node}} = 1 < \text{Threshold} = 1.5$$

$$\text{Disp}(\text{Eraser}_4) = \frac{1}{1}$$

$$\text{Disp}(\text{Eraser}_2) = \frac{3}{3}$$

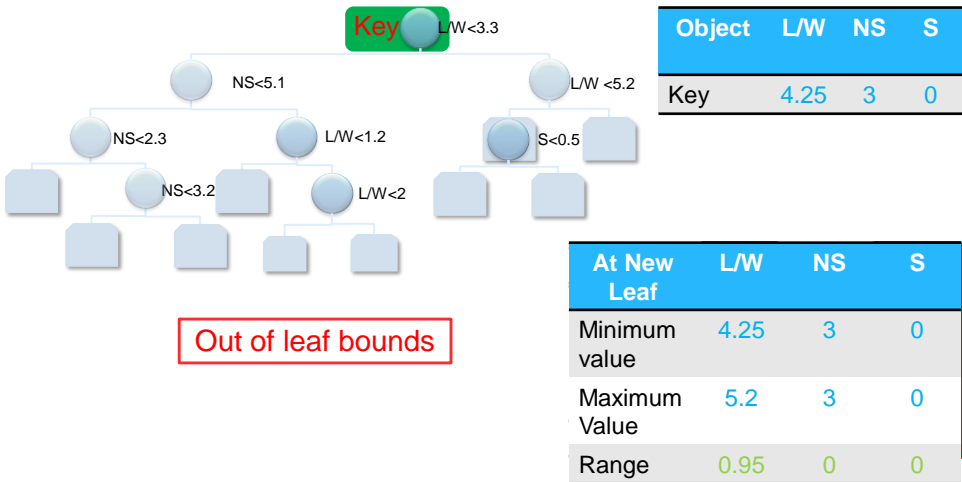
$$\text{Disp}(\text{Eraser}_3) = \frac{1}{2}$$

$$\text{Disp}(\text{Eraser}_1) = \frac{2}{6}$$



RRCF Anomalies: Example


- Consider the new object *Key*, by sliding the window by one point





RRCF Anomalies: Example

Fall 2019



$$\text{Displacement}(\text{Data Point}_{\text{depth}}) = \frac{\text{leaves of sibling node}}{\text{leaves of current node}}$$

Key \rightarrow **Outlier**
 Should discard changes and Keep Original Tree

$$\text{Disp}(\text{Key}_3) = \frac{1}{2} \quad \text{Disp}(\text{Key}_2) = \frac{1}{1} \quad \text{Disp}(\text{Key}_1) = \frac{6}{3}$$

$$\text{CoDisp}(\text{Eraser}) = \max(\text{Displacement}) = 2 > \text{Threshold} = 1.5$$



RRCF Complexity

Fall 2019

RRCF Building Complexity

RRCF Updating Complexity

$$O(t(n))$$

$$O(t(\log_2(n)))$$

- Worst case : perfect binary tree, where each data point is a leaf
- Where,
 - t is the number of trees
 - n is the maximum number of data points in a tree
 - $(2n - 1)$ is the total number of nodes of a perfect binary tree with n leaves
 - $\log_2(n)$ is the minimal height of a perfect binary tree with n leaves
 - Tree height is increased by 1 after processing a new data point
- Worst case : perfect binary tree, where each new data point is a new leaf at the maximal height updating the entire subtree



63



RRCF Summary


- RRCF chooses a **splitting feature proportionally to its value range** and not uniformly
 - Formally proves that tree **structure updates** are equivalent as if the data point were available during training
 - Tackles the presence of **irrelevant dimensions** for outlier isolation
- RRCF **determines distance based anomalies** without focusing on the specifics of the distance function
 - Define **anomalies** using **high average of collusive displacement** instead of low average tree height
 - Tackles finer differences (**masking**) between outliers and inliers
- RRCF hyper-parameters include the CoDisp **threshold**
 - As **CoDisp is not normalized**, we should first normalize the maximal tree displacements before computing their average in the forest
 - It is open how to tune threshold, in the normalized range $[0, 1]$, to accurately separate normal from abnormal points
- RRCF is implemented in the AWS Data Analytics Engine



Half Space Trees (HST)

[Tan et al, IJCAI 2011; Chen et al, MLJ 2015]

- Randomly **construct** small binary trees as **iForest**, but
 - For each feature define its **workspace**
 - A HST bisects data space using randomly selected features and the **middle value** of its workspace
 - A node in a HST captures the number of points within a particular subspace, called **Mass** profile
- Incrementally **update** the **Mass** profiles of the tree/s, for every new **tumbling window**, to learn the density of the stream
 - **Transfer** the non-zero **mass** of the last window to the **model** computed in the previous window
- **Anomalies** can be inferred using the current window points
 - Points are **ranked** in their ascending **score** order
 - The **lower** the **score** of a point, the more **probable** to be an **anomaly**




Example: A HST on a 2D Stream



HST Stream Snapshot: Example

- Given a stream of objects, lets say 8 objects



- Process the stream using tumbling windows of length $w = 4$
 - Start by collecting the first w objects, as a **Reference** window

Object	L/W	NS	S
Box	1	6	1
Block	1.6	6	1
Bead	1	2	1
Penny	1	3	1



HST Workspace: Example

- Compute the Workspace of each feature, using the **reference** window
- A. **Work range:** Minimum, Maximum, Split and Range values of features

Object	L/W	NS	S
Box	1	6	1
Block	1.6	6	1
Bead	1	2	1
Penny	1	3	1

(A)

Workrange	L/W	NS	S
Minimum	1	2	1
Maximum	1.6	6	1
Split	1.55	5.38	1
Range	1.10	6.76	0

- B. **Work space:** Max, Min values of the features' work range

(B)

$$\text{Max} = \text{Split} + \text{Range}$$

$$\text{Min} = \text{Split} - \text{Range}$$

Workspace	L/W	NS	S
Max	2.66	12.14	1
Min	0.45	-1.38	1



Fall 2019

HST Construction: Example

- Construct a HST using the features' work space


Workspace	L/W	NS	S
Max	2.66	12.14	1
Min	0.45	-1.38	1

- Randomly select a feature F
- Split data space using the mid value SP_F
- Initialize the mass profile

- Control tree growth, using the objects of the reference window

Object	L/W	NS	S
Box	1	6	1
Block	1.6	6	1
Bead	1	2	1
Penny	1	3	1

- Activated: Max Height Limit = 2
- Activated: Object Size Limit = 1



68




Fall 2019

HST Mass Update: Example

- Update the Mass_r Profile of the HST, using the objects of the reference window
 - Increase by one the r mass of visited nodes

Object	L/W	NS	S
Box	1	6	1
Block	1.6	6	1
Bead	1	2	1
Penny	1	3	1




69



HST Stream Snapshot: Example

- Process of the reference window has been **completed**



- Continue by collecting the **next** w objects, as a **Latest** window

Object	L/W	NS	S
Screw	8	3	0
Dime	1	3	1
knob	1	4	1
Battery	5	3	1




HST Anomaly Detection: Example

Detect **anomalies** in the latest window, using the **current** HST Model

Object	L/W	NS	S
Screw	8	3	0
Dime	1	3	1
knob	1	4	1
Battery	5	3	1

Object	Score	Rank
Screw	2	1 st
Dime	8	2 nd
knob	8	2 nd
Battery	2	1 st

1
0
0
1



$$\text{Predict}_{\text{Obj}} = \text{Score}_{\text{Obj}} < \text{mean}(\text{score}) : 1 ? 0$$








HST Stream Snapshot: Example

- Process of the current latest window has been **completed**



- Continue by collecting the **final** w objects as a **Latest** window

Object	L/W	NS	S
Easer	1.6	6	1
Key	4.25	3	0
Paper	10	2	1
Pin	1.14	3	0



HST Anomaly Detection: Example

- Detect **anomalies** in the latest window, using the current HST Model

Object	L/W	NS	S
Easer	1.6	6	1
Key	4.25	3	0
Paper	10	2	1
Pin	1.14	3	0


Object	Score	Rank
Easer	6	1 st
Key	6	1 st
Paper	6	1 st
Pin	16	2 nd

1

1

1

0



$$\text{Predict}_{\text{Obj}} = \text{Score}_{\text{obj}} < \text{mean}(\text{score}) : 1 \text{ ? } 0$$





HST Complexity

HST Building Complexity

HST Updating Complexity

$$O(t * 2^{h+1})$$

$$O(t * h * w)$$

- Worst case : perfect binary tree, where each data point is a leaf
- Where,
 - t is the number of trees
 - h is the max height of a tree
 - w is the number of points in a window
 - $(2^{h+1}-1)$ is the number of all nodes of a perfect binary tree of max height h
- Complexities are **constant** (amortized) when the h , t and w are set



HST Summary

- HST **workspace** estimates the feature value ranges of future data in a stream, resulting to more **accurate** mass updates in stable distributions
- HST **structure** initialized using the first *reference* window, and **remains fixed** for all streaming data
 - Incrementally **updating only** the mass profiles of the HST model
- HST **ranks anomalies** using a **scoring function** that considers **only the mass profiles of the leaf nodes**
- HST fails to support concept drift when the **statistical characteristics of features evolve over time**
 - **incorrectly** selected **features** may lead to **inaccurate mass updates**
- One of the HST hyper-parameters is the **size** of tumbling windows
 - How we can tune window size to tackle concept drift/shift?




Qualitative Comparison

Algorithm	Novelty	Splitting Feature	Splitting Value	Window Type	Anomaly Score	Model Update	Anomaly Types
IF	Avoids computation of distances/densities	Uniform	Uniform	χ	Tree height	batch	Global & Local
RRCF	Addresses Masking & Irrelevant Features	Proportional to the feature range	Uniform	Slide	Collusive Disp. (sensitive to neighborhood)	Incremental tree updates	Global & Local; <i>Sample</i> based range subspace
HST	High Speed Data Streams	Uniform	Mid Value	Tumble	Mass Profiles (sensitive to neighborhood)	Mini-batch update of mass profiles	Global & Local; fixed range subspace



Questions?





Ensemble Learning Techniques

Fall 2019

Can we turn a weak learner into a strong learner?

	Bagging (Bootstrap Aggregating)	Boosting	Stacking
Data Partitioning	Random samples are drawn with replacement	every new subsets contains the samples that were (likely to be) misclassified by previous models	Various
Goal	Minimize Variance	Increase predictive power	Both
Exploiting methods	Random Subspace	Gradient descent	Logistic Regression
Fusion of models	(Weighted) Average	(Weighted) Majority Vote	Meta model to estimate the weights



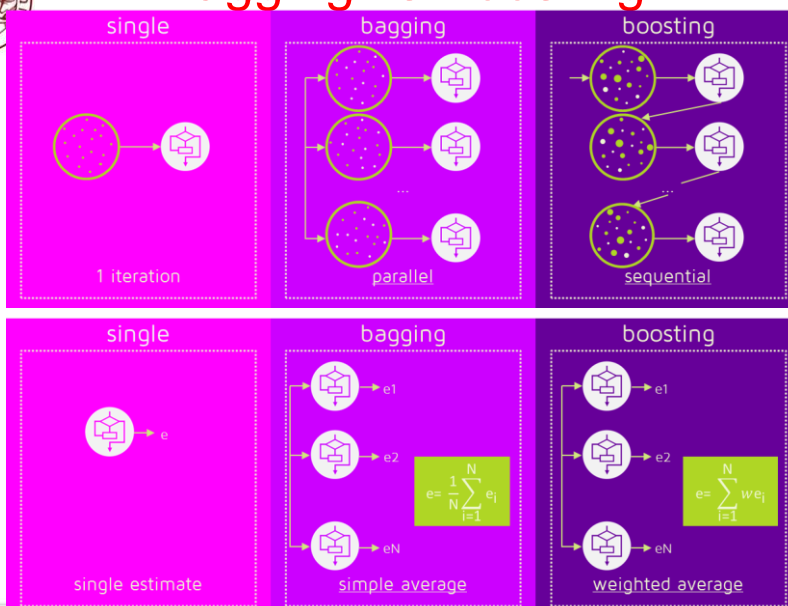
<https://stats.stackexchange.com/questions/18891/bagging-boosting-and-stacking-in-machine-learning>

82



Bagging vs Boosting

Fall 2019



<https://www.kdnuggets.com/2017/11/difference-bagging-boosting.html>

83



Bibliography

Fall 2019

- C. Aggarwal, S. Sathe. Outlier Ensembles: An Introduction (1st ed.). 2017 Springer.
- C. Aggarwal. *Outlier Analysis*. 2013 Springer.
- E. Schubert. Generalized and efficient outlier detection for spatial, temporal, & high-dimensional data mining. Ph.D. LMU München, 2013
- M. Facure. Semi-Supervised Anomaly Detection Survey
<https://www.kaggle.com/quietrags/semi-supervised-anomaly-detection-survey>
- R. Domingues, M. Filippone, P. Michiardi, J. Zouaoui. A comparative evaluation of outlier detection algorithms: Experiments and analyses. *Pattern Recognition* Volume 74, Feb. 2018
- A. Zimek, E. Schubert, and H.-P. Kriegel. A survey on unsupervised outlier detection in high-dimensional numerical data, *Statistical Analysis and Data Mining*, 5(5), 2012
- A. Muallem, S. Shetty, J. Pan, J. Zhao, B. Biswal. Hoeffding Tree Algorithms for Anomaly Detection in Streaming Datasets: A Survey. *JIS*. Vol.8 No.4, October 2017
- A. Lavin, S. Ahmad. Evaluating Real-time Anomaly Detection Algorithms – The Numenta Anomaly Benchmark, *ICMLA*, December 2015.
- Fei Tony Liu, Kai Ming Ting, and Zhi-Hua Zhou. Isolation-Based Anomaly Detection. *ACM Trans. Knowl. Discov. Data* 6, 1, Article 3 March 2012, 39 pages.
- Z. Ding, M. Fei. An anomaly detection approach based on isolation forest algorithm for streaming data using sliding window, *IFAC Proceedings Volumes* 46 (2013) 12–17.
- S. C. Tan, K. M. Ting, T. F. Liu. Fast anomaly detection for streaming data. In *Proc. of the IJCAI 2011 Volume Two*, AAAI Press 1511-1516.
- K. M. Ting, J. R. Wells. Multi-dimensional Mass Estimation and Mass-based Clustering. In *Proc. of the ICDM 2010*, IEEE, Washington, DC, USA, 511-520.
- T. Bandaragoda, K. Ting, D. Albrecht, F. Liu, Y. Zhu, J. R. Wells. Isolation-based anomaly detection using nearest-neighbor ensembles. *Int. Journal Comput. Intelligence*, Wiley 2018



Bibliography

Fall 2019

- J. Marcus Hughes, Anomaly Detection and Robust Random Cut Trees, 2019 [Github](#)
- S. Guha, N. Mishra, G. Roy, O. Schrijvers. Robust random cut forest based anomaly detection on streams. In *Proc. ICML 2016*, Vol. 48. [JMLR.org](http://jmlr.org) 2712-2721.
- A. Siddiqui, A. Fern, T. Dietterich, W.-K. Wong. Sequential Feature Explanations for Anomaly Detection. *Workshop on Outlier Definition, Detection, and Description (ODDx3)*, August 2015
- B. Chen, K.M. Ting, T. Washio, G. Haffari. Half-Space Mass: A maximally robust and efficient data depth method. *Machine Learning*. 100(2-3), 2015, 677-699.
- K. M. Ting, G.-T. Zhou, F. T. L., J. S. C. Tan. Mass estimation and its applications. In *Proc. KDD 2010*. ACM, New York, NY, USA, 989-998.
- K. Wu, K. Zhang, W. Fan, A. Edwards, P. S. Yu. RS-Forest: A Rapid Density Estimator for Streaming Anomaly Detection. *IEEE ICDM 2014*, pp 600–609.
- J.R. Wells, K.M. Ting, T. Washio. LiNearN: A New Approach to Nearest Neighbour Density Estimator. *Pattern Recognition*. Vol.47, 8, 2014, 2702-2720.
- B. Schölkopf, R. Williamson, A. Smola, J. Shawe-Taylor, J. Platt. Support vector method for novelty detection. In *Proc. NIPS 1999*, MIT Press, Cambridge.
- **Distributed, Data Stream Processing**
 - M. Dias de Assuno, A. da Silva Veith, R. Buyya. Distributed data stream processing and edge computing. *J. Netw. Comput. Appl.* 103, C, February 2018, 1-17.
 - A. Shukla, Y. Simmhan. Benchmarking Distributed Stream Processing Platforms for IoT Applications. In *Proc. of the Technology Conference on Performance Evaluation and Benchmarking* February 2017, LNCS.
 - N. Tatbul, S. Zdonik, J. Meehan, C. Aslantas, M. Stonebraker, K. Tufte, C. Giossi, H. Quach. Handling Shared, Mutable State in Stream Processing with Correctness Guarantees. *IEEE Data Eng. Bull.* 38(4): 94-104 2015



Anomaly Detection Software Libraries

Fall 2019

- Collection of anomaly detection examples : https://github.com/shubhomoydas/ad_examples
- Isolation Forest (iForest) implementation
 - in R: <https://r-forge.r-project.org/projects/iforest/>, <https://rdrr.io/rforge/IsolationForest/>, <https://sourceforge.net/projects/iforest/>, <https://github.com/Zelazny7/isoform>, <https://github.com/zmzhang/IOS>
 - in sklearn: <http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.IsolationForest.html>
 - In python: https://github.com/mgckind/iso_forest
 - In spark: <https://github.com/titicaca/spark-iforest>
 - In Java: <https://github.com/bnjmn/weka/blob/master/packages/internal/isolationForest/src/main/java/weka/classifiers/misc/IsolationForest.java>
 - in Go: <https://github.com/e-XpertSolutions/go-iforest>
- One Class SVM (1CSVM) implementation
 - in R: <https://gumroad.com/l/nbjri> (download the supplemental zip file at http://univprofblog.html.xdomain.jp/code/R_scripts_functions.zip)
 - In Python: <https://gum.co/oPLZ> (download the supplemental zip file at <http://univprofblog.html.xdomain.jp/code/supportingfunctions.zip>)
 - In Java: <https://github.com/jnioche/libsvm-java>
 - In MS Azure: <https://docs.microsoft.com/en-us/azure/machine-learning/studio-module-reference/one-class-support-vector-machine>
- Unsupervised Anomaly detection algorithms
 - In Rapid Miner: <http://madm.dfki.de/rapidminer/anomalydetection>
- Anomaly Detection using One-Class Neural Networks: <https://github.com/raghavchalapathy/oc-nn>



86



License

Fall 2019

- These slides are made available under a Creative Commons Attribution-ShareAlike license (CC BY-SA 3.0): <http://creativecommons.org/licenses/by-sa/3.0/>



- You can share and remix this work, provided that you keep the attribution to the original authors intact, and that, if you alter, transform, or build upon this work, you may distribute the resulting work only under the same or similar license to this one



87